

Class 7: Machine Learning 1

Wade Ingersoll (PID: 69038080)


Today we will begin our exploration of some “classical” machine learning approaches. We will start with clustering.

Let’s first make up some data to cluster where we know what the answer should be.

```
rnorm(10)
```

```
[1] 0.956021743 -0.355340256 1.105906724 0.009855852 -1.069203832  
[6] -0.260899404 -0.376616831 0.068803141 -1.200704482 0.105846924
```

```
hist( rnorm(1000) )
```



```

x <- c( rnorm(30, mean=-3), rnorm(30, mean=3) )
y <- rev(x)


x <- cbind(x, y)
head(x)

```

	x	y
[1,]	-4.845343	3.722339
[2,]	-4.652340	3.732024
[3,]	-1.168317	1.778395
[4,]	-2.446093	3.930176
[5,]	-4.096390	4.494303
[6,]	-3.341280	2.809756

A look at x with `plot()`

```
plot(x)
```



Then main function in “base” R for K-means clustering is called `kmeans()`

```
k <- kmeans(x, centers = 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1  3.239057 -3.143580
2 -3.143580  3.239057

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 58.60648 58.60648  
(between_SS / total_SS = 91.2 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. How big are the clusters (i.e. their size)?

k\$size


[1] 30 30

Q. What clusters do my data points reside in?

k\$cluster

Q. Make a plot of our data colored by cluster assignment - i.e. Make a result figure...

```
plot(x, col=k$cluster)  
points(k$centers, col="blue", pch=15)
```




Q. Cluster with k-means into 4 clusters and plot the results as above.

```
x <- c( rnorm(30, mean=-3), rnorm(30, mean=3) )
y <- rev(x)
x <- cbind(x, y)

k <- kmeans(x, centers = 4)

plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15)
```



Q. Run kmeans with center (i.e. values of k) equal 1 to 6

```
k1 <- kmeans(x, centers=1)$tot.withinss
k2 <- kmeans(x, centers=2)$tot.withinss
k3 <- kmeans(x, centers=3)$tot.withinss
k4 <- kmeans(x, centers=4)$tot.withinss
k5 <- kmeans(x, centers=5)$tot.withinss
k6 <- kmeans(x, centers=6)$tot.withinss

answer <- c(k1, k2, k3, k4, k5, k6)
```


Or use a for loop

```
answer <- NULL
for (i in 1:6) {
  answer <- c(answer, kmeans(x, centers=i)$tot.withinss)
}
answer
```

```
[1] 1141.72973   91.17920   75.68312   66.70166   59.77674   44.88282
```

Make a “scree-plot”

```
plot(answer, typ="b")
```



Hierarchical Clustering

The main function in “base” R for this is called `hclust()`


```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)
abline(h=7, col="red")
```

Cluster Dendrogram




```
d  
hclust (*, "complete")
```

To obtain clusters from our `hclust()` result object `hc` we “cut” the tree to yield different sub branches. For this we use the `cutree()` function

```
cutree(hc, k=7)
```

```
[1] 1 2 3 3 3 3 3 3 1 1 3 3 3 3 3 3 1 3 3 3 2 3 2 3 1 3 1 2 3 3 4 4 5 6 4 6 4 5  
[39] 4 5 4 4 4 6 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 5 7
```

```
library(pheatmap)  
pheatmap(x)
```



Principal Component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1: How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

A: Using the code below, I found there are 17 rows and 5 columns

```
# Complete the following code to find out how many rows and columns are in x?
dim(x)
```

```
[1] 17 5
```

Preview the first 6 rows

```
head(x)
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139

```
# Note how the minus indexing works
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

		England	Wales	Scotland	N.Ireland
Cheese		105	103	103	66
Carcass_meat		245	227	242	267
Other_meat		685	803	750	586
Fish		147	160	122	93
Fats_and_oils		193	235	184	209
Sugars		156	175	147	139

Q2: Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

A: The code below can solve this problem. This version works better if you plan to run the code multiple times.

```
x <- read.csv(url, row.names=1)
head(x)
```

		England	Wales	Scotland	N.Ireland
Cheese		105	103	103	66
Carcass_meat		245	227	242	267
Other_meat		685	803	750	586
Fish		147	160	122	93
Fats_and_oils		193	235	184	209
Sugars		156	175	147	139


Spotting major differences and trends

```
rainbow(5)
```

```
[1] "#FF0000" "#CCFF00" "#00FF66" "#0066FF" "#CC00FF"
```

Using base R

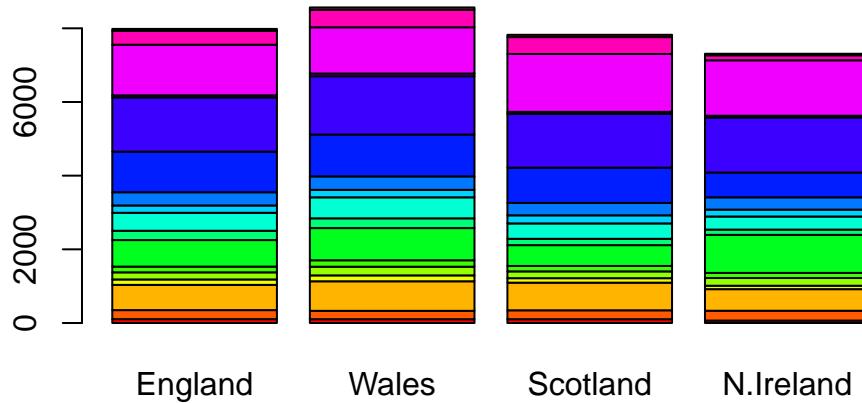
```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

A: Changing `beside` to `F`

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Convert data to long format for ggplot with pivot_longer()

```
library(tidyr)

x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
               names_to = "Country",
               values_to = "Consumption")

dim(x_long)
```

[1] 68 3

Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure?


A: Adding an argument `position = "stack"` within `geom_col()`

```

library(ggplot2)

ggplot(x_long) +
  aes(x = Country, y = Consumption, fill = Food) +
  geom_col(position = "stack") +
  theme_bw()

```



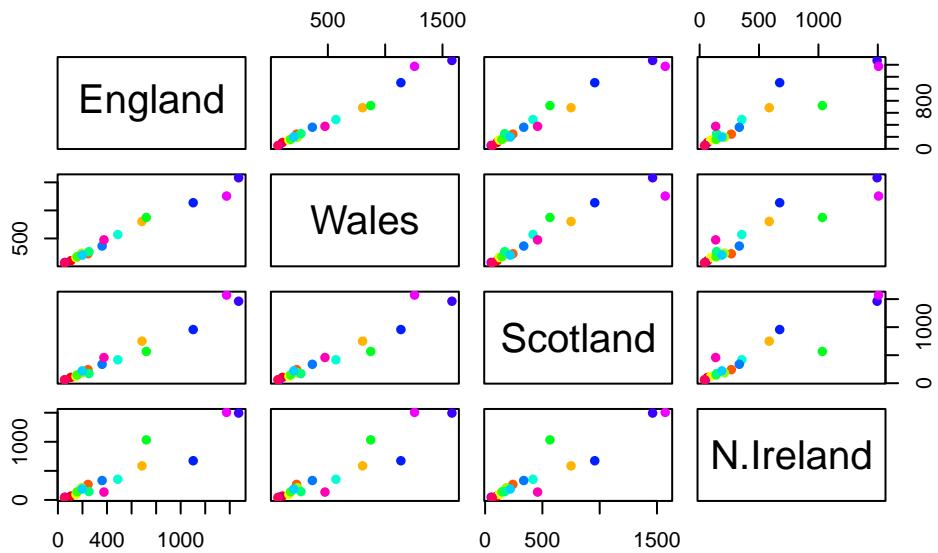
Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

A: Each point is a food and when they lie on the diagonal, it means the two countries eat it at the same rate


```

pairs(x, col=rainbow(nrow(x)), pch=16)

```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

A: It looks like Wales and England are quite similar in their consumption of these foods. It is still quite difficult to tell what is going on in the dataset.

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`.

As we want to do PCA on the food data for the different countries, we will want the foods in the columns.

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

Our result object is called `pca` and it has a `$x` component that we will look at first

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.


```

library(ggplot2)

# Create a data frame for plotting
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

# Plot PC1 vs PC2 with ggplot
ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()

```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.


```

library(ggplot2)

cols <- c("orange", "red", "blue", "darkgreen")


```

```
# Plot PC1 vs PC2 with ggplot
ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, col=cols) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



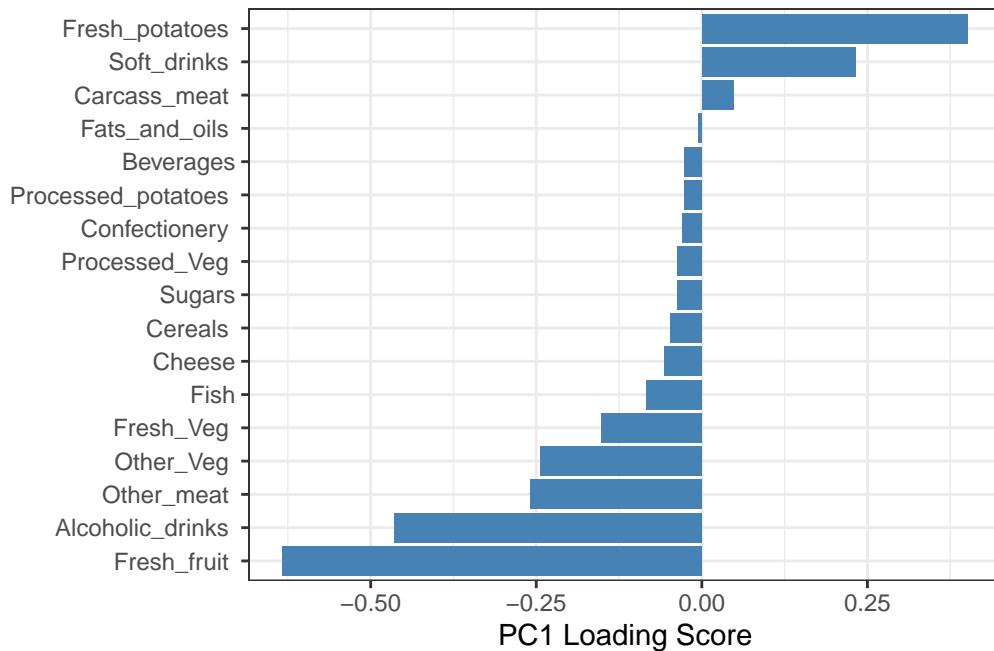
Another major result out of the PCA is the so-called “variable loadings” or \$rotation that tells us how the original variables (foods) contribute to the PCs (i.e. our new axis).

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```




Lets focus on PC1 as it accounts for > 90% of variance

```
ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
ggplot(pca$rotation) +
  aes(x = PC2,
      y = reorder(rownames(pca$rotation), PC2)) +
  geom_col(fill = "skyblue") +
  xlab("PC2 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set? How many PCs do you think it will take to have a useful overview of this data set (see below)?

A: Using `dim()`, there are 100 genes and 10 samples. I think it will take just 1 PC to have a useful overview of the data set.

```

dim(rna.data)


[1] 100 10

## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

# Create data frame for plotting
df <- as.data.frame(pca$x)
df$Sample <- rownames(df)

## Plot with ggplot
ggplot(df) +
  aes(x = PC1, y = PC2, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, size = 3) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()

```



Examine variation in original data

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251
	PC8	PC9	PC10				
Standard deviation	0.62065	0.60342	3.344e-15				
Proportion of Variance	0.00385	0.00364	0.000e+00				
Cumulative Proportion	0.99636	1.00000	1.000e+00				


A quick scree plot summary of this Proportion of Variance for each PC can be obtained using ggplot:

```
# Calculate variance explained
pca.var <- pca$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)

# Create scree plot data
scree_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  Variance = pca.var[1:10]
)

ggplot(scree_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  ggtitle("Quick scree plot") +
  xlab("Principal Component") +
  ylab("Variance") +
  theme_bw()
```

Quick scree plot



Percent variance is often more informative to look at

```
pca.var.per
```

```
[1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

Create percent variance scree plot


```
scree_pct_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  PercentVariation = pca.var.per[1:10]
)

ggplot(scree_pct_df) +
  aes(x = PC, y = PercentVariation) +
  geom_col(fill = "steelblue") +
  ggtitle("Scree Plot") +
  xlab("Principal Component") +
```

```

ylab("Percent Variation") +
theme_bw()

```



Now lets make our main PCA plot a bit more attractive and useful...

```

## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

# Add condition to data frame
df$condition <- substr(df$Sample, 1, 2)
df$color <- colvec


ggplot(df) +
  aes(x = PC1, y = PC2, color = color, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, hjust = 0.5, show.legend = FALSE) +
  scale_color_identity() +
  xlab(paste0("PC1 (", pca.var.per[1], "%)")) +

```

```

ylab(paste0("PC2 (", pca.var.per[2], "%)")) +
theme_bw()

```



Optional Gene Loadings

For demonstration purposes let's find the top 10 measurements (genes) that contribute most to pc1 in either direction (+ or -).

```

loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes

```

```

[1] "gene100" "gene66"   "gene45"   "gene68"   "gene98"   "gene60"   "gene21"
[8] "gene56"   "gene10"   "gene90"

```