

Class 19: Investigating Pertussis

Wade Ingersoll (PID: A69038080)

Table of contents

Background	1
The CMI-PB project	4
Side-Note: Working with dates	6
Joining multiple tables	10
4. Examine IgG Ab titer levels	13

Background

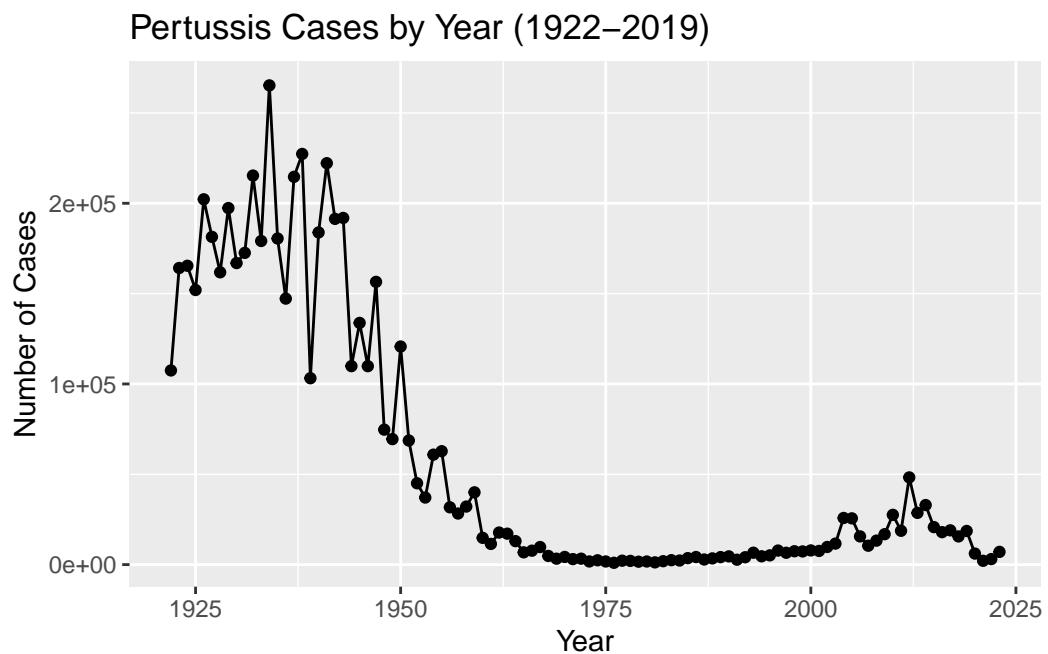
Pertussis (a.k.a. Whooping cough) is a highly contagious lung infection caused by the bacteria *B. pertussis*.

The CDC tracks case numbers in the US and makes this data available online

> **Q1:** With the help of the R “addin” package **datapasta** assign the CDC pertussis case number data to a data frame called **cdc** and use **ggplot** to make a plot of cases numbers over time.

```
library(ggplot2)

ggplot(cdc) +
  aes(year, cases) +
  geom_point() +
  geom_line() +
  labs(x="Year", y="Number of Cases", title = "Pertussis Cases by Year (1922-2019)")
```



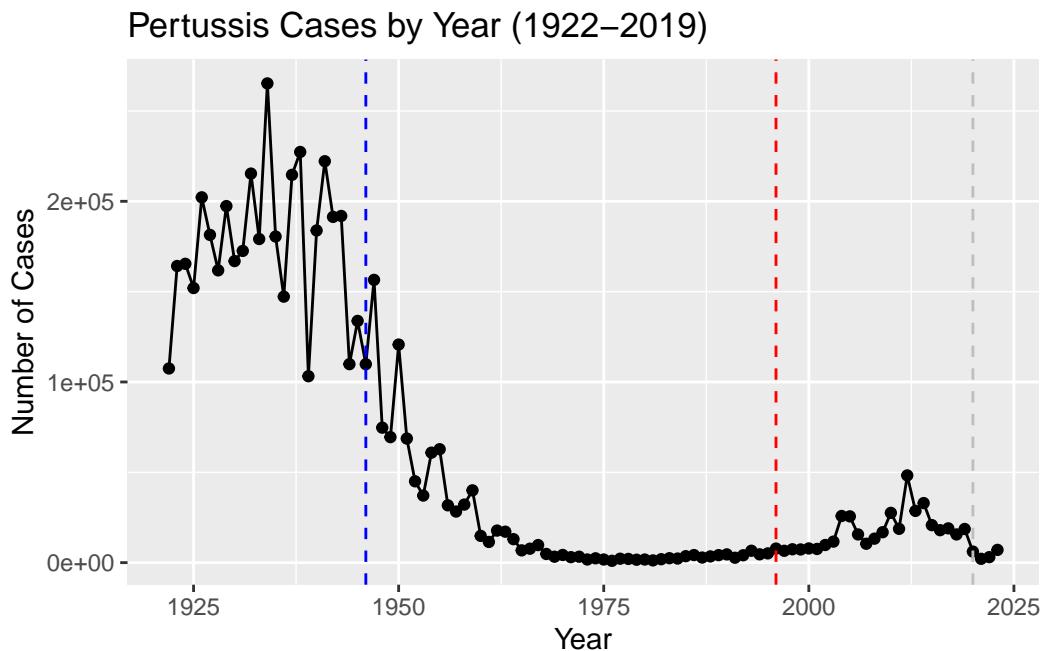
Q2: Using the ggplot `geom_vline()` function add lines to your previous plot for the 1946 introduction of the wP vaccine and the 1996 switch to aP vaccine (see example in the hint below). What do you notice? Finally, add a line for 2020.

Answer: I noticed that approximately 10 years after switching to the aP vaccine, cases began to rise.

Q2.1: Describe what happened after the introduction of the aP vaccine? Do you have a possible explanation for the observed trend?

Answer: After introduction of the aP vaccine, the number of cases increased. This could be due to aP's decreased ability to confer lasting immunity compared to wP.

```
ggplot(cdc) +  
  aes(year, cases) +  
  geom_point() +  
  geom_line() +  
  geom_vline(xintercept = 1946, col="blue", linetype="dashed") +  
  geom_vline(xintercept = 1996, col="red", linetype="dashed") +  
  geom_vline(xintercept = 2020, col="grey", linetype="dashed") +  
  labs(x="Year", y="Number of Cases", title = "Pertussis Cases by Year (1922-2019)")
```



The CMI-PB project

The CMI-Pertussis Boost (PB) project focuses on gathering data on this very topic. What is distinct between aP and wP individuals over time when they encounter Pertussis again.

They make their data available via a JSON format returning API. We can read JSON format with the `read_json()` function from the `jsonlite` package with `install.packages("jsonlite")`

```
library(jsonlite)
subject <- read_json("http://cmi-pb.org/api/v5_1/subject", simplifyVector = TRUE)
head(subject)
```

	subject_id	infancy_vac	biological_sex	ethnicity	race
1	1	wP	Female	Not Hispanic or Latino	White
2	2	wP	Female	Not Hispanic or Latino	White
3	3	wP	Female		Unknown
4	4	wP		Male Not Hispanic or Latino	Asian
5	5	wP		Male Not Hispanic or Latino	Asian
6	6	wP	Female	Not Hispanic or Latino	White

	year_of_birth	date_of_boost	dataset
1	1986-01-01	2016-09-12	2020_dataset
2	1968-01-01	2019-01-28	2020_dataset
3	1983-01-01	2016-10-10	2020_dataset
4	1988-01-01	2016-08-29	2020_dataset
5	1991-01-01	2016-08-29	2020_dataset
6	1988-01-01	2016-10-10	2020_dataset

Q3: How many “subjects” (or individuals) are in this dataset?

Answer: 172 (see code below for calculations)

```
nrow(subject)
```

```
[1] 172
```

Q4: How many aP and wP infancy vaccinated subjects are in the dataset?

Answer: There are 87 aP and 85 wP (see code below for calculations)

```
table(subject$infancy_vac)
```

```
aP wP  
87 85
```

Q5: How many Male and Female subjects/patients are in the dataset?

Answer: There are 112 females and 60 males (see code below for calculations)

```
table(subject$biological_sex)
```

```
Female   Male  
112      60
```

Q6: What is the breakdown of race and biological sex (e.g. number of Asian females, White males etc...)?

Answer: See code below for breakdown

```
table(subject$race, subject$biological_sex)
```

	Female	Male
American Indian/Alaska Native	0	1
Asian	32	12
Black or African American	2	3
More Than One Race	15	4
Native Hawaiian or Other Pacific Islander	1	1
Unknown or Not Reported	14	7
White	48	32

Side-Note: Working with dates

Two of the columns of `subject` contain dates in the Year-Month-Day format. Recall from our last mini-project that dates and times can be annoying to work with at the best of times. However, in R we have the excellent lubridate package, which can make life allot easier. Here is a quick example to get you started:

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

date, intersect, setdiff, union

What is today's date (at the time I am writing this obviously)

```
today()
```

[1] "2025-12-06"

How many days have passed since new year 2000

```
today() - ymd("2000-01-01")
```

Time difference of 9471 days

What is this in years?

```
time_length( today() - ymd("2000-01-01") , "years")
```

[1] 25.93018

Note that here we are using the `ymd()` function to tell lubridate the format of our particular date and then the `time_length()` function to convert days to years.

Q7: Using this approach determine (i) the average age of wP individuals, (ii) the average age of aP individuals; and (iii) are they significantly different?

Answer: The average age of **wP = 37** and **aP = 28**. These values are statistically significant with a **p-value of 2.372101e-23**. See code below for calculations.

```
# Using ymd() function on the appropriate column of our subject data frame together with today's date  
subject$age <- today() - ymd(subject$year_of_birth)  
  
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
# Next we can use dplyr's filter() function to limit ourselves to a particular subset of subjects  
ap <- subject %>% filter(infancy_vac == "aP")  
  
round( summary( time_length( ap$age, "years" ) ) )
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
23	27	28	28	29	35

```
# wP  
wp <- subject %>% filter(infancy_vac == "wP")  
round( summary( time_length( wp$age, "years" ) ) )
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
23	33	35	37	40	58

```
# Are they significantly different?
t <- t.test(ap$age, wp$age)
t$p.value
```

Time difference of 2.372101e-23 days

Q8: Determine the age of all individuals at time of boost?

Answer: See code below for calculations

```
int <- ymd(subject$date_of_boost) - ymd(subject$year_of_birth)
age_at_boost <- time_length(int, "year")
subject$age_at_boost <- age_at_boost
head(subject)
```

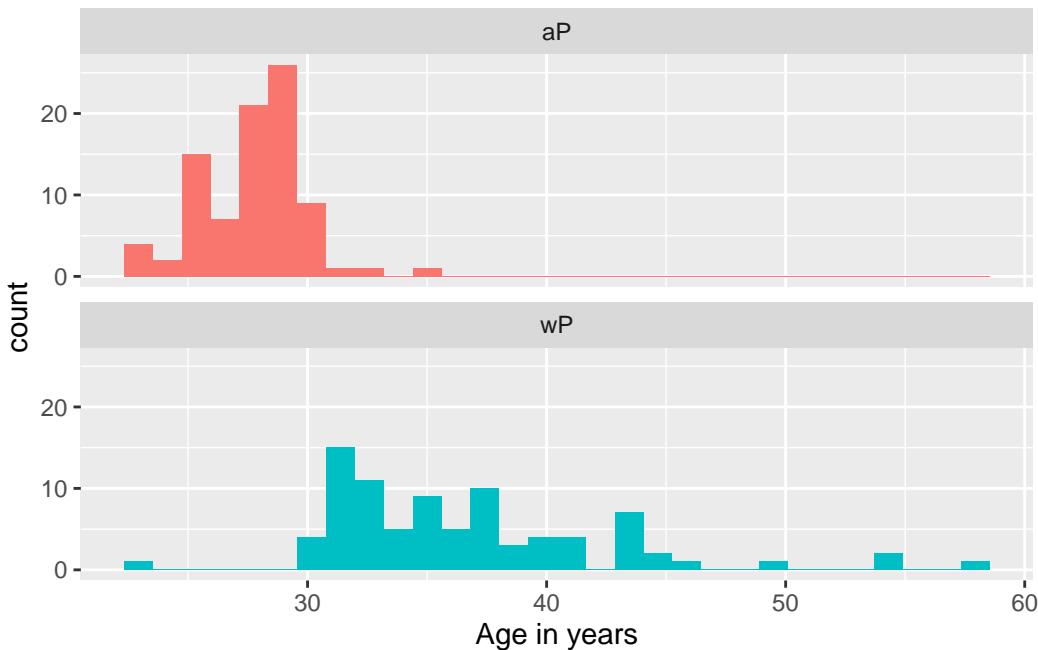
	subject_id	infancy_vac	biological_sex	ethnicity	race
1	1	wP	Female	Not Hispanic or Latino	White
2	2	wP	Female	Not Hispanic or Latino	White
3	3	wP	Female		Unknown White
4	4	wP		Male Not Hispanic or Latino	Asian
5	5	wP		Male Not Hispanic or Latino	Asian
6	6	wP	Female	Not Hispanic or Latino	White
	year_of_birth	date_of_boost	dataset	age	age_at_boost
1	1986-01-01	2016-09-12	2020_dataset	14584 days	30.69678
2	1968-01-01	2019-01-28	2020_dataset	21159 days	51.07461
3	1983-01-01	2016-10-10	2020_dataset	15680 days	33.77413
4	1988-01-01	2016-08-29	2020_dataset	13854 days	28.65982
5	1991-01-01	2016-08-29	2020_dataset	12758 days	25.65914
6	1988-01-01	2016-10-10	2020_dataset	13854 days	28.77481

Q9a: With the help of a faceted boxplot or histogram (see below), do you think these two groups are significantly different?

Answer: These groups definitely seem different. This observation is supported by a t-test p-value of **2.372101e-23**

```
ggplot(subject) +  
  aes(time_length(age, "year"),  
      fill=as.factor(infancy_vac)) +  
  geom_histogram(show.legend=FALSE) +  
  facet_wrap(vars(infancy_vac), nrow=2) +  
  xlab("Age in years")
```

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.



```
x <- t.test(time_length( wp$age, "years" ),  
            time_length( ap$age, "years" ))
```

```
x$p.value
```

```
[1] 2.372101e-23
```

Joining multiple tables

Read the specimen and ab_titer tables into R and store the data as `specimen` and `titer` named data frames.

```
# Complete the API URLs...
specimen <- read_json("http://cmi-pb.org/api/v5_1/specimen", simplifyVector = TRUE)
ab_titer <- read_json("http://cmi-pb.org/api/v5_1/plasma_ab_titer", simplifyVector = TRUE)
```

To know whether a given `specimen_id` comes from an aP or wP individual we need to link (a.k.a. “join” or merge) our `specimen` and `subject` data frames. The excellent `dplyr` package (that we have used previously) has a family of `join()` functions that can help us with this common task:

Q9b: Complete the code to join specimen and subject tables to make a new merged data frame containing all specimen records along with their associated subject details:

```
meta <- inner_join(specimen, subject)
```

Joining with `by = join_by(subject_id)`

```
dim(meta)
```

```
[1] 1503   15
```

```
head(meta)
```

	specimen_id	subject_id	actual_day_relative_to_boost	planned_day_relative_to_boost	specimen_type	visit	infancy_vac	biological_sex
1	1	1	-3	0	Blood	1	wP	Female
2	2	1	1	1	Blood	2	wP	Female
3	3	1	3	3	Blood	3	wP	Female
4	4	1	7	7	Blood	4	wP	Female
5	5	1	11					
6	6	1	32					

5		14	Blood	5	wP	Female
6		30	Blood	6	wP	Female
		ethnicity	race	year_of_birth	date_of_boost	dataset
1	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset	
2	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset	
3	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset	
4	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset	
5	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset	
6	Not Hispanic or Latino	White	1986-01-01	2016-09-12	2020_dataset	
		age	age_at_boost			
1	14584 days	30.69678				
2	14584 days	30.69678				
3	14584 days	30.69678				
4	14584 days	30.69678				
5	14584 days	30.69678				
6	14584 days	30.69678				

Q10: Now using the same procedure join meta with titer data so we can further analyze this data in terms of visit aP/wP, male/female etc.

```
abdata <- inner_join(ab_titer, meta)
```

Joining with `by = join_by(specimen_id)`

```
dim(abdata)
```

```
[1] 61956    22
```

Q11: How many specimens (i.e. entries in abdata) do we have for each isotype?

Answer: See code below for answer

```
table(abdata$isotype)
```

IgE	IgG	IgG1	IgG2	IgG3	IgG4
6698	7265	11993	12000	12000	12000

Q12: What are the different \$dataset values in abdata and what do you notice about the number of rows for the most “recent” dataset?

Answer: For the most recent dataset (2023), the number of rows increased compared to the two previous years. See code below for other \$dataset values.

```
table(abdata$dataset)
```

2020_dataset	2021_dataset	2022_dataset	2023_dataset
31520	8085	7301	15050

4. Examine IgG Ab titer levels

Now using our joined/merged/linked abdata dataset filter() for IgG isotype.

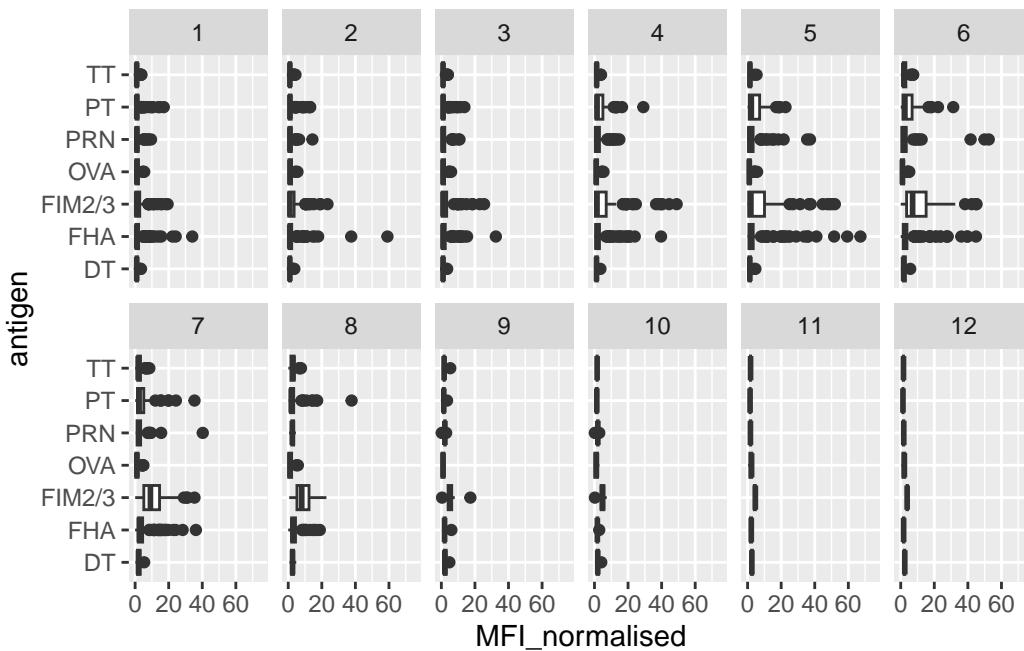
```
igg <- abdata %>% filter(isotype == "IgG")
head(igg)
```

	specimen_id	isotype	is_antigen_specific	antigen	MFI	MFI_normalised
1	1	IgG	TRUE	PT	68.56614	3.736992
2	1	IgG	TRUE	PRN	332.12718	2.602350
3	1	IgG	TRUE	FHA	1887.12263	34.050956
4	19	IgG	TRUE	PT	20.11607	1.096366
5	19	IgG	TRUE	PRN	976.67419	7.652635
6	19	IgG	TRUE	FHA	60.76626	1.096457
			unit	lower_limit_of_detection	subject_id	actual_day_relative_to_boost
1	IU/ML			0.530000	1	-3
2	IU/ML			6.205949	1	-3
3	IU/ML			4.679535	1	-3
4	IU/ML			0.530000	3	-3
5	IU/ML			6.205949	3	-3
6	IU/ML			4.679535	3	-3
			planned_day_relative_to_boost	specimen_type	visit	infancy_vac biological_sex
1			0	Blood	1	wP Female
2			0	Blood	1	wP Female
3			0	Blood	1	wP Female
4			0	Blood	1	wP Female
5			0	Blood	1	wP Female
6			0	Blood	1	wP Female
			ethnicity	race	year_of_birth	date_of_boost dataset
1	Not Hispanic or Latino	White		1986-01-01	2016-09-12	2020_dataset
2	Not Hispanic or Latino	White		1986-01-01	2016-09-12	2020_dataset
3	Not Hispanic or Latino	White		1986-01-01	2016-09-12	2020_dataset
4		Unknown	White	1983-01-01	2016-10-10	2020_dataset
5		Unknown	White	1983-01-01	2016-10-10	2020_dataset
6		Unknown	White	1983-01-01	2016-10-10	2020_dataset
			age	age_at_boost		
1	14584	days		30.69678		
2	14584	days		30.69678		
3	14584	days		30.69678		
4	15680	days		33.77413		
5	15680	days		33.77413		
6	15680	days		33.77413		

Q13: Complete the following code to make a summary boxplot of Ab titer levels (MFI) for all antigens:

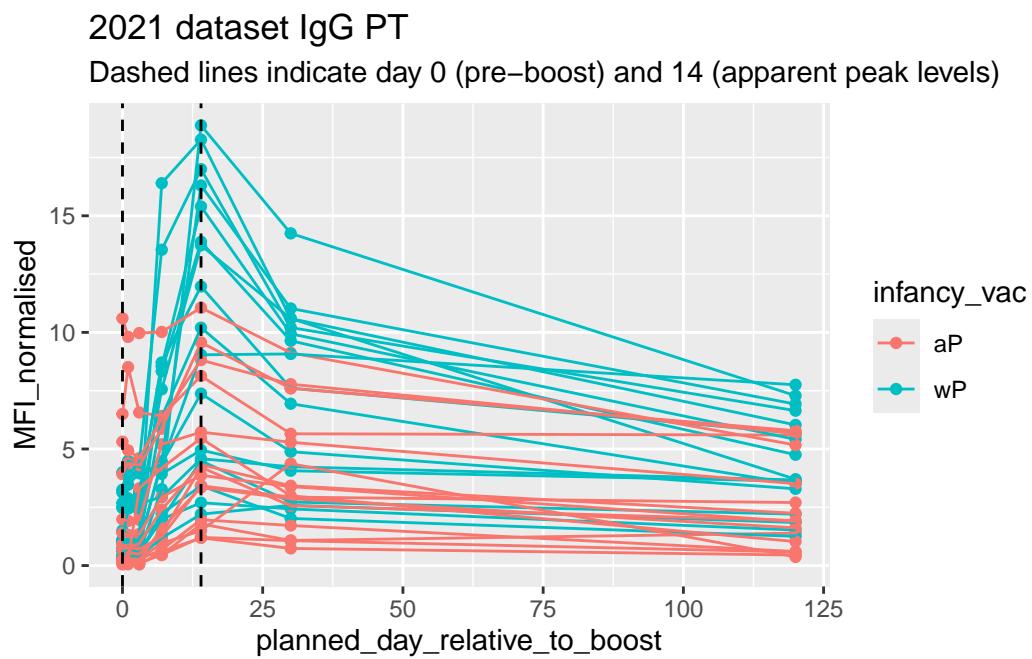
```
ggplot(igg) +  
  aes(MFI_normalised, antigen) +  
  geom_boxplot() +  
  xlim(0,75) +  
  facet_wrap(vars(visit), nrow=2)
```

Warning: Removed 5 rows containing non-finite outside the scale range
(`stat_boxplot()`).



Q14: Antigen levels time-course plot

```
abdata.21 <- abdata %>% filter(dataset == "2021_dataset")  
  
abdata.21 %>%  
  filter(isotype == "IgG", antigen == "PT") %>%  
  ggplot() +  
    aes(x=planned_day_relative_to_boost,  
        y=MFI_normalised,  
        col=infancy_vac,  
        group=subject_id) +  
    geom_point() +  
    geom_line() +  
    geom_vline(xintercept=0, linetype="dashed") +  
    geom_vline(xintercept=14, linetype="dashed") +  
    labs(title="2021 dataset IgG PT",  
         subtitle = "Dashed lines indicate day 0 (pre-boost) and 14 (apparent peak levels)")
```



Q14 continued

```
igg <- igg %>% filter(visit %in% 1:8)
ggplot(igg) +
  aes(MFI_normalised, antigen, col=infancy_vac) +
  geom_boxplot() +
  facet_wrap(~visit, nrow=2)
```

